

Programming 2




Figure 1: Episode duration vs episodes

Architecture:

- Linear layer (in_dim=4, out_dim=64, activation=Relu)
- Linear layer (in_dim=64, out_dim=2, activation=None)

Hyperparameters:

- Batch size = 128
- GAMMA = 0.999
- EPS_START = 1.0
- EPS_END = 0.04
- EPS_DECAY = 0.0095
- TARGET_UPDATE = 10
- MEMORY_CAPACITY = 10000
- Optimizer: Adam
- learning rate: 1e-3
- num episodes = 100
- Network weight initialization: Xavier normal

Performance on test episodes: Total episodes = 10

Episode duration: = [291,267,216,291,269,298,271,303,381,223]

Mean duration for testing= 281.0

Link for video: Test episodes video